

Computational Learning Theory

Agam Goyal

January 28, 2023

Lec 01: Introduction

Computational learning theory studies machine learning from a theoretical computer science point of view. **Machine Learning:** Programs that enable automatic extraction of useful information from “raw” data, and which improve their performance via interaction with this data. Examples: Classification, Clustering, Prediction, etc.

Overview of Computational Learning Theory

A Learning Model should specify the following:

- (1) “Who” is learning?
 - A computer program, usually restricted
 - (i) Polynomial run-time
 - (ii) Small sample complexity
 - Particular format for output hypothesis (for the interpretability of the model)
- (2) “What” are we learning?
 - Skills and environment
 - “Classification” rule: Boolean function
- (3) “How” does the learner get information?
 - Learner is given examples as point-label pairs: $(x, f(x))$
 - (i) “Passive” learning - only observes the information provided by the environment without influencing it
 - (ii) Assumptions about X :
 - Randomly chosen
 - Maliciously chosen (Adversary models)
 - Provided by teacher (Adversarial teacher model – worst-case scenario)
 - Learner makes queries
 - (i) “Active” learning - interacts with the environment at training time, by posing queries or performing experiments

Is the data ever noisy or incomplete?

The data is usually almost always noisy and/or incomplete, and it is useful to study the robustness of learning models to these datasets.

What “prior rule” does the learner have?

The typical assumption is that there is an a priori representational scheme for the function being learned known as the **target concept**.

Performance Criterion

- How to measure accuracy?
- Online vs Offline/Batch Learning Methods

Lecture 02: Online Mistake-Bound learning Models

Setup

\mathbf{X} : instance space

\mathcal{C} : concept (a Boolean function on \mathbf{X}) — $c : \mathbf{X} \rightarrow \{0, 1\}$

\mathcal{C} : Concept Class (set of all Boolean functions) — $c \in \mathcal{C}$

Typically, $\mathbf{X} = \{0, 1\}^n$ or $\mathbf{X} = \mathbb{R}^n$ and the learner knows \mathcal{C} and \mathbf{X} , but doesn't know the individual concept target c which is a specific function.

Example #1

$$\mathbf{X} = \{0, 1\}^n$$

$\mathcal{C} =$ All **monotone conjunctions** over x_1, \dots, x_n

A monotone conjunction is an AND of variables such that no variable in the conjunction is negated.

$$|\mathcal{C}| = 2^n \rightarrow \text{All possible subsets}$$

Example of a concept: $c(\mathbf{x}) = x_3 \wedge x_5 \wedge x_6$

Example #2

$$\mathbf{X} = \{0, 1\}^n$$

$\mathcal{C} =$ All conjunctions over x_1, \dots, x_n

$$|\mathcal{C}| = 3^n \rightarrow \text{Each literal can exist as itself, its conjugate, or not exist}$$

Example of a concept: $c(\mathbf{x}) = \bar{x}_2 \wedge \bar{x}_4 \wedge x_7 \wedge \bar{x}_8$

Example #3

$$\mathbf{X} = \{0, 1\}^n$$

$\mathcal{C} =$ All **DNFs** over x_1, \dots, x_n with $\leq n^2$ terms

A Disjunctive Normal Form (DNF) is an OR of ANDs of literals.

$$|\mathcal{C}| = \text{Number of terms in the DNF formula}$$

Example of a concept: $c(\mathbf{x}) = (x_2 x_3) \vee (x_1 \bar{x}_2 x_4) \vee (\bar{x}_2 \bar{x}_4 \bar{x}_5 x_7) \vee x_7 \bar{x}_9$

We can extend this idea to a **k-DNF**: a DNF with $\leq k$ variables.

Example #4

$$\mathbf{X} = \mathbb{R}^n$$

$\mathcal{C} =$ All Linear Threshold Functions (LTFs / Halfspaces)

Example of a concept: $c : \mathbb{R}^n \rightarrow \{0, 1\}$

$$\exists n \text{ real weights } w_1, w_2, \dots, w_n \in \mathbb{R} \text{ and } \theta \in \mathbb{R}, \text{ s.t. } \forall \mathbf{x} \in \mathbb{R}^n, c(\mathbf{x}) = \begin{cases} 1, & \text{if } \langle \mathbf{w}, \mathbf{x} \rangle \geq \theta \\ 0, & \text{otherwise} \end{cases}$$

where \mathbf{w} is the vector of weights w_1, w_2, \dots, w_n , and $\mathbf{x} \in \mathbf{X}$.

Online Mistake-Bound Learning Models

Algorithm 1 Online Mistake-Bound Learning Model

```
1: A learning session proceeds in a sequence of trials
2: Throughout the session, learner maintains hypothesis  $h : \mathbf{X} \rightarrow \{0, 1\}$ 
3: while In a trial do:
4:   Learner is given some  $\mathbf{x} \in \mathbf{X}$ 
5:   Learner outputs  $h(\mathbf{x})$  as prediction
6:   Learner is told the true  $c(\mathbf{x})$ 
7:   if  $h(\mathbf{x}) \neq c(\mathbf{x})$  then
8:     Learner is charged for the mistake
9:   Learner may update  $h$ 
10:
```

Performance Measure = Number of Mistakes

Definition: Mistake Bound (M)

A learning algorithm A has a Mistake Bound M for a concept class \mathcal{C} if

$(\forall c \in \mathcal{C}) (\forall \text{ sequence of examples from } \mathbf{X}) \quad A \text{ makes } \leq M \text{ mistakes.}$

Observations:

- If \mathbf{X} is finite, the learner can always achieve $M = |\mathbf{X}|$.
- If \mathcal{C} is finite, the learner can always achieve $M = |\mathcal{C}| - 1$.

Example #1

$\mathbf{X} = \{0, 1, \dots, 2^n - 1\}$

$\mathcal{C} =$ All initial intervals $[0, \alpha)$ where $\alpha \in \mathbb{Z}^+ \cup 0$

$|\mathcal{C}| = 2^n + 1 \rightarrow$ Includes every possible integer from 0 to $2^n - 1$ and also the interval that even excludes 0.

Consider the scenario where we want to split the points in \mathbf{X} into 0s and 1s or $\mathbf{X} \rightarrow \{0, 1\}$. Everything part of the open interval is assigned a 0 and everything else is assigned a 1. Our goal is to find an initial interval which clearly separates these two classes.

Note that we can use **binary search** to learn \mathcal{C} with $M = \log_2(2^n) \Rightarrow \boxed{M = n}$. Our initial hypothesis h in this case would just be the initial half interval $\{0, 1, \dots, 2^{n-1}\}$. And then if the learner makes a mistake for some input \mathbf{x} , we can update the interval accordingly. So essentially, each mistake eliminates \geq Half of the remaining possibilities.

Example #2

$\mathbf{X} = [0, 1]$

$\mathcal{C} =$ All initial intervals $[0, \alpha)$ where $\alpha \in [0, 1]$

In this case, the mistake bound $\boxed{M = \infty}$ because of examples arbitrarily close to the boundaries we choose. So we may never really get to the *exact* answer we desire. However, in reality we would often just want to achieve a result that would be a *good enough* approximation.

Lecture 03: Online Learning of Disjunctions and Decision Lists